


Huawei Summer Internship: Enabling SVE/SME on triton-cpu


Daniel Antonio Martínez Sánchez

Compiler Engineer Intern (OSPP)
Huawei

September 24, 2025

1. Introduction
2. Challenges
3. Current status

- Enable the new SVE and SME ARM vector extensions into the open euler triton-cpu 



+ SVE/SME

So, how do we do that?

- We can just hardcode the intrinsics for relevant ops 😞
- But we have some really nice SVE/SME work already done
 - The team at Cambridge already has really good performance with SVE/SME
 - Why don't just use that in triton?
 - That's a challenge because triton is not really easy to handle in this aspect

The issue with triton

Triton Uses it's own IR almost entirely

TTIR

```
%8 = tt.addptr %7, %4 : tensor<64x!tt.ptr<f32>>, tensor<64xi32>  
%9 = tt.load %8, %6: tensor<64x!tt.ptr<f32>>
```

TTCIR

```
%6 = tt.addptr %arg0, %1 : !tt.ptr<f32>, i32  
%7 = triton_cpu.ptr_to_memref %6 : <f32> -> memref<128xf32>  
%8 = vector.maskedload %7[%c0], %5, %cst -> vector<128xf32>
```

LLVM IR

```
%19 = getelementptr float, ptr %0, i64 %18  
%20 = call <128xfloat> @llvm.masked.load.v128f32(ptr %19, 4, <128xi1>  
%17)
```


X86: AVX512 instructions: `vmovups`, `vaddps %zmm`

The issue with triton

- But on Cambridge they use standard MLIR...
- is there any way to make it work?



- Triton shared is a middle layer between triton and MLIR
- It translates from the triton dialects to standard MLIR
- This enables us to still use triton but have all of the advantages of MLIR!



```
tt.func @kernel(%afloat : !tt.ptr<bf16>, %res : !tt.ptr<bf16>) {  
  %0 = tt.make_range {end = 128 : i32, start = 0 : i32} : tensor<128xi32>  
  %1 = tt.splat %afloat : (!tt.ptr<bf16>) -> tensor<128x!tt.ptr<bf16>>  
  %2 = tt.addptr %1, %0 : tensor<128x!tt.ptr<bf16>>, tensor<128xi32>  
  %afm = tt.load %2 : tensor<128x!tt.ptr<bf16>>  
  %3 = "tt.reduce"(%afm) ({  
    ^bb0(%arg5: bf16, %arg6: bf16):  
      %21 = arith.addf %arg5, %arg6 : bf16  
      tt.reduce.return %21 : bf16  
  }) {axis = 0 : i32} : (tensor<128xbf16>) -> bf16  
  tt.store %res, %3 : !tt.ptr<bf16>  
  tt.return  
}
```



```
func.func @kernel(%arg0: memref<*xbf16>, %arg1: memref<*xbf16>, %arg2: i32, %arg3: i32, %arg4: i32) {  
    %cst = arith.constant 0.000000e+00 : f32  
    %reinterpret_cast = memref.reinterpret_cast %arg0 to offset: [0], sizes: [128], strides: [1] :  
        memref<*xbf16> to memref<128xbf16, strided<[1]>>  
    %alloc = memref.alloc() : memref<128xbf16>  
    memref.copy %reinterpret_cast, %alloc : memref<128xbf16, strided<[1]>> to memref<128xbf16>  
    %0 = bufferization.to_tensor %alloc restrict writable : memref<128xbf16>  
    %1 = bufferization.alloc_tensor() : tensor<f32>  
    %inserted = tensor.insert %cst into %1[] : tensor<f32>  
    %reduced = linalg.reduce ins(%0 : tensor<128xbf16>) outs(%inserted : tensor<f32>) dimensions = [0]  
    (%in: bf16, %init: f32) {  
        %3 = arith.extf %in : bf16 to f32  
        %4 = arith.addf %3, %init : f32  
        linalg.yield %4 : f32  
    }  
    %extracted = tensor.extract %reduced[] : tensor<f32>  
    %2 = arith.truncf %extracted : f32 to bf16  
    %reinterpret_cast_0 = memref.reinterpret_cast %arg1 to offset: [0], sizes: [1], strides: [1] :  
        memref<*xbf16> to memref<1xbf16, strided<[1]>>  
    affine.store %2, %reinterpret_cast_0[0] : memref<1xbf16, strided<[1]>>  
    return  
}
```

How good is triton-shared?

- Performance
 - Triton itself is a DSL designed for GPU's !
 - Performance can get really compromised depending on the source triton kernel.
- Reliability
 - Triton-shared is an experimental project 🧪
 - It's under active development

Example of bad code:

```
for m in range(0, M, BLOCK_SIZE_M):  
    # Do in parallel  
    for n in range(0, N, BLOCK_SIZE_N):  
        acc = zeros((BLOCK_SIZE_M, BLOCK_SIZE_N), dtype=float32)  
        for k in range(0, K, BLOCK_SIZE_K):  
            a = A[m : m+BLOCK_SIZE_M, k : k+BLOCK_SIZE_K]  
            b = B[k : k+BLOCK_SIZE_K, n : n+BLOCK_SIZE_N]  
            acc += dot(a, b)  
        C[m : m+BLOCK_SIZE_M, n : n+BLOCK_SIZE_N] = acc
```

Example of performant code: ✓

```
def bare_matmul(X, Y, Z, M, N, K, BLOCK_SIZE: tl.constexpr):
    pid_x = tl.program_id(0)  # block row id
    pid_y = tl.program_id(1)  # block column id

    offs_x = pid_x * BLOCK_SIZE + tl.arange(0, BLOCK_SIZE)
    offs_y = pid_y * BLOCK_SIZE + tl.arange(0, BLOCK_SIZE)

    x = tl.load(X + offs_x[:, None] * K + offs_y[None, :])
    y = tl.load(Y + offs_x[:, None] * N + offs_y[None, :])

    z = tl.dot(x, y)

    tl.store(Z + offs_x[:, None] * N + offs_y[None, :], z)
```

At the start of the project we had ~5000 failing test (out of 6658) 😱:

- Unsupported datatypes (fp8)
- Unsupported triton operations
- Lots of LLVM bugs

A quick summary of the new operations supported:


- Reductions with complex bodies (`tt.reduce`)
- Scans (`tt.scan`)
- Bitcasts (`tt.bitcast`)
- Atomic read, modify, write (`tt.atomic_rmw`)
- Batched matmul (`tt.dot` with 3d inputs)

And a quick summary of LLVM bugs and patches

- Crash when running SME on transforms dialect
- Add support for f8e5m2 & f8e4m3fn in arith expands-ops
- Linalg missing RecursiveMemoryEffects
- Crash causes by liveness analysis in remove-dead-values
- New SVE/SME pipelines


With all of those fixes and additions we have just...

```
==== 96 failed, 6562 passed ====
```




96 failing test left .



What about the SME/SVE optimized pipelines?

- Pipelines had to be open-sourced from the Cambridge team 
- SME pipeline was open sourced early and his already ported to triton
- SVE pipeline was open sourced recently and there is still work todo

Future work

- Finish porting SVE pipeline 
- Fix remaining failing test 
- Benchmark performance 

I wrote some documentation about this:

- Project log here
- Reproduction instructions here